

512 Price Discovery in Inference Markets

Why Not All AI Tokens Are Equal

Commodities Are Priced by Constraints, Not Labels

In mature commodity markets, units are nominally identical — but **economically distinct**.

A kilowatt-hour is always one kilowatt-hour.

Yet prices vary dramatically based on:

- time
- location
- reliability
- source
- compliance guarantees

Inference tokens behave identically.

What Actually Determines the Price of an Inference Token

An inference token clears at a price defined by its **constraint envelope**, not by its nominal compute size.

Key pricing dimensions include:

1. Latency Class

- sub-5 ms (control-critical)
- sub-20 ms (interactive)
- sub-100 ms (batch-tolerant)

Lower latency tokens command structural premiums.

This is physics, not preference.

2. Hardware Class

- consumer GPU
- enterprise GPU
- specialized accelerator

Different failure modes, reliability, and throughput profiles

→ different clearing prices.

3. Energy Source & ESG Profile

Just as power markets differentiate:

- coal
- gas
- nuclear
- renewable

Inference tokens will be priced by:

- energy intensity
- carbon footprint
- heat reuse
- jurisdictional energy mix

For regulated buyers, **ESG-qualified inference** is not optional.

4. Jurisdiction & Data Sovereignty

Tokens executed under:

- EU data residency
- financial compliance zones
- defense-grade isolation

...carry embedded legal guarantees.

Those guarantees have cost.

And therefore price.

5. SLA & Finality Guarantees

Some tokens will include:

- uptime guarantees
- execution attestations
- irreversible settlement

Others will be best-effort.

Just like:

- firm power vs interruptible power
-

Market Stratification Is Inevitable

As in electricity markets:

- high-value tokens do not disappear
- low-value tokens are not useless

They clear in **different markets**.

This produces:

- premium inference
- spot inference
- interruptible inference
- green inference
- sovereign inference

No central planner chooses this.

Markets discover it.

Why Price Discovery Forces Standardization

Price discovery only works if:

- units are comparable
- attributes are disclosed
- settlement is final

This is why electricity required:

- standardized metering
- grid codes
- disclosure rules

Inference markets require the same.

512 provides:

- mandatory disclosure of constraints

- non-negotiable consent rules
- auditable finality
- universal exit

Without these, price signals collapse.

The Quiet CFO Insight

CFOs do not care about AI ideology.

They care about:

- cost curves
- predictability
- auditability
- compliance risk

Inference tokens with:

- known constraints
- known provenance
- known settlement

...are **financeable assets**.

Unconstrained inference is not.

Locking Statement

A kilowatt is never “just” a kilowatt.

An inference token is never “just” compute.

Price emerges from constraints — and constraints require rules.

512 does not set prices.

It makes prices possible.